

Compulsory 2

Sigbjørn Fjelland

2022-05-06

Task 1

Problem 1.1

```
library(bootstrap)
library(boot)
library(ggplot2)
library(data.table)
library(Rcpp)
library(VGAM)
```

```
## Loading required package: stats4
## Loading required package: splines
##
## Attaching package: 'VGAM'
## The following objects are masked from 'package:boot':
##
##      logit, simplex
## The following object is masked from 'package:bootstrap':
##
##      hormone
```

```
library(MASS)
```

```
# Presenting the data
data <- scor
data.open <- c("mec", "vec")
data.closed <- c("alg", "ana", "sta")
all.sets <- c(data.open, data.closed)

#constructing groups that are to be correlated as function for the boot package

correlating.groups <- function(data, ind){
  mek <- data[ind, 1]
  vec <- data[ind, 2]
  alg <- data[ind, 3]
  ana <- data[ind, 4]
  sta <- data[ind, 5]

  c(cor(mek,vec), cor(alg,ana), cor(alg, sta), cor(ana,sta))
}
```

```
}
```

```
boot(data = data, statistic = correlating.groups, R = 200)
```

```
##
```

```
## ORDINARY NONPARAMETRIC BOOTSTRAP
```

```
##
```

```
##
```

```
## Call:
```

```
## boot(data = data, statistic = correlating.groups, R = 200)
```

```
##
```

```
##
```

```
## Bootstrap Statistics :
```

```
##      original      bias    std. error
```

```
## t1* 0.5534052  0.0015970223  0.07349188
```

```
## t2* 0.7108059  0.0032297071  0.05060925
```

```
## t3* 0.6647357 -0.0008492849  0.05792702
```

```
## t4* 0.6071743  0.0096818818  0.06458413
```

problem 1.2

```
theta.hat <- function(data, ind){
  mek <- data[ind, 1]
  vec <- data[ind, 2]
  alg <- data[ind, 3]
  ana <- data[ind, 4]
  sta <- data[ind, 5]

  data <- cbind(mek, vec, alg, ana, sta)
  lambda <- eigen(cov(data))$values

  lambda[1]/sum(lambda)
}

theta.boot <- boot(data = data, statistic = theta.hat, R=200, mle = eigen(cov(data)))
theta.boot

##
## ORDINARY NONPARAMETRIC BOOTSTRAP
##
##
## Call:
## boot(data = data, statistic = theta.hat, R = 200, mle = eigen(cov(data)))
##
##
## Bootstrap Statistics :
##      original      bias    std. error
## t1* 0.619115 0.004075202  0.0474656
```

problem 1.3

```
print(boot.ci(theta.boot,type = "perc"))
```

```
## BOOTSTRAP CONFIDENCE INTERVAL CALCULATIONS
## Based on 200 bootstrap replicates
##
## CALL :
## boot.ci(boot.out = theta.boot, type = "perc")
##
## Intervals :
## Level      Percentile
## 95%      ( 0.5181,  0.7093 )
## Calculations and Intervals on Original Scale
## Some percentile intervals may be unstable
```

Task 2

problem 2.1 MH-Sampler with rayleigh distribution as proposal distribution

```
set.seed(123)

#Proposal distribution (Rayleigh)
f <- function(x, sigma) {
  if (any(x < 0)) return (0)
  stopifnot(sigma > 0)
  return((x / sigma^2) * exp(-x^2 / (2*sigma^2)))
}

Nsim <- 20000 #total simulations
burnin = 5000 #sample to burn
sigma <- 4

x <- numeric(Nsim)
x[1] <- rgamma(1, rate = 1, shape = 1) # start value to simplify index in loop
k <- 0 ## conter for rejects ##
u <- runif(Nsim)

for (i in 2:Nsim) {
  xt <- x[i-1]
  y <- rgamma(1, shape = xt, rate = 1)
  num <- f(y, sigma) * dgamma(xt, rate = 1, shape = y) #the numerator of r(xt,y)
  den <- f(xt, sigma) * dgamma(y, rate = 1, shape = xt) #the denominator of r(xt,y)
  if (u[i] <= num/den) x[i] <- y else {
    x[i] <- xt
    k <- k+1 # Number of y being rejected
  }
}


index<-seq(10,Nsim,by=10) ## take index every 10 steps

#sampled distribution from MH (and library for compartion)
sampled_ray_dist <- x[burnin:Nsim]
raylib <- rrayleigh(Nsim+1-burnin, scale = sigma)

#quantile function
a <- ppoints(1000)
QR <- sigma * sqrt(-2 * log(1 - a))


#Hists
hist(sampled_ray_dist, breaks="scott",
      main="MH-sampler and QF of Raylight", xlab="x", freq=FALSE)
lines(QR, f(QR, 4))
```

MH-sampler and QF of Raylight



```
hist (raylib,freq=FALSE,  
      main = "Rayl. dist. from lib. sampler and QF.", xlab = "x")  
lines(QR, f(QR, 4))
```

Rayl. dist. from lib. sampler and QF.



Problem 2.2

Vi Benytter normal fordelingen som symetrisk distribusjon, med $\sigma = (0.05, 0.5, 2, 16)$, og på gøy lager vi en sidekick med laplace

```
rw.Metropolis.normal <- function(n, sigma, x0, N) {
  x <- numeric(N)
  x[1] <- x0
  u <- runif(N)
  k <- 0
  for (i in 2:N) {
    y <- rnorm(1, x[i-1], sigma) #rlaplace(n=1, location=x[i-1], scale = sigma) #
    if (u[i] <= (dt(y, n) / dt(x[i-1], n)))
      x[i] <- y else {
        x[i] <- x[i-1]
        k <- k + 1
      }
  }
  return(list(x=x, k=k))
}

rw.Metropolis.laplace <- function(n, sigma, x0, N) {
  x <- numeric(N)
  x[1] <- x0
  u <- runif(N)
  k <- 0
  for (i in 2:N) {
    y <- rlaplace(n=1, location=x[i-1], scale = sigma) #
    if (u[i] <= (dt(y, n) / dt(x[i-1], n)))
      x[i] <- y else {
        x[i] <- x[i-1]
        k <- k + 1
      }
  }
  return(list(x=x, k=k))
}
```

Random Walk for samplet fra normal fordelingen:

```
n <- 4 #degrees of freedom for target Student t dist.
N <- 2000
sigma <- c(.05, .5, 2, 16)
x0 <- 25
rw1.n <- rw.Metropolis.normal(n, sigma[1], x0, N)
rw2.n <- rw.Metropolis.normal(n, sigma[2], x0, N)
rw3.n <- rw.Metropolis.normal(n, sigma[3], x0, N)
rw4.n <- rw.Metropolis.normal(n, sigma[4], x0, N)

# Result from the random walks given sigma vector
print(c(rw1.n$k, rw2.n$k, rw3.n$k, rw4.n$k))

## [1] 7 225 911 1804

#Proportion of rejections within [0.15, 0.5]
print(c((2000 - length(which(rw1.n$x > 0.5)) - length(which(rw1.n$x < 0.15)))/N,
        (2000 - length(which(rw2.n$x > 0.5)) - length(which(rw2.n$x < 0.15)))/N,
        (2000 - length(which(rw3.n$x > 0.5)) - length(which(rw3.n$x < 0.15)))/N,
        (2000 - length(which(rw4.n$x > 0.5)) - length(which(rw4.n$x < 0.15)))/N))

## [1] 0.000 0.099 0.116 0.167
```


Vi ser at både $\sigma = 2$ og $\sigma = 16$ har en lik andelen innenfor intervallet. For å få en illustrasjon på hvilke σ som konvergerer plotter vi resultatet:

```
#Plot of sigma1 and sigma2
par(mfrow=c(1,2))


plot(rw1.n$x, type = "l", main="sigma = 0.05")
abline(h=c(0.15,0.5))

plot(rw2.n$x, type = "l",main="sigma = 0.5")
abline(h=c(0.15,0.5))
```


sigma = 0.05




sigma = 0.5




```
#Plot of sigma1 and sigma2
par(mfrow=c(1,2))
plot(rw3.n$x, type = "l",main="sigma = 2")
abline(h=c(0.15,0.5))

plot(rw4.n$x, type = "l",main="sigma = 16")
abline(h=c(0.15,0.5))
```

sigma = 2



sigma = 16



En ser på plottet at både $\sigma = 0.5$, $\sigma = 2$ og $\sigma = 16$ konvergerer innenfor intervallet, men $\sigma = 0.5$ trenger en lengre periode med burning før den setter seg i en omegn av intervallet. $\sigma = 16$ ser noe mindre effektiv ut. Foretrukken σ å optimalisere rundt ville vert $\sigma = 2$.

og så til vår lille sidekick:

```
n <- 4 #degrees of freedom for target Student t dist.
N <- 2000
sigma <- c(.05, .5, 2, 16)
x0 <- 25
rw1.l <- rw.Metropolis.laplace(n, sigma[1], x0, N)
rw2.l <- rw.Metropolis.laplace(n, sigma[2], x0, N)
rw3.l <- rw.Metropolis.laplace(n, sigma[3], x0, N)
rw4.l <- rw.Metropolis.laplace(n, sigma[4], x0, N)

# Result from the random walks given sigma vector
print(c(rw1.l$k, rw2.l$k, rw3.l$k, rw4.l$k))

## [1] 9 327 953 1751

#Proportion of rejections within [0.15, 0.5]
print(c((2000 - length(which(rw1.l$x > 0.5)) - length(which(rw1.l$x < 0.15)))/N,
        (2000 - length(which(rw2.l$x > 0.5)) - length(which(rw2.l$x < 0.15)))/N,
        (2000 - length(which(rw3.l$x > 0.5)) - length(which(rw3.l$x < 0.15)))/N,
        (2000 - length(which(rw4.l$x > 0.5)) - length(which(rw4.l$x < 0.15)))/N))



## [1] 0.000 0.108 0.119 0.216

For å se for hvilke  $\sigma$  som konvergerer

#Plot of sigma1 and sigma2
par(mfrow=c(1,2))
```

```
plot(rw1.l$x, type = "l")
abline(h=c(0.15,0.5))
```



```
plot(rw2.l$x, type = "l")
abline(h=c(0.15,0.5))
```



```
#Plot of sigma1 and sigma2
```

```
par(mfrow=c(1,2))
plot(rw3.l$x, type = "l")
abline(h=c(0.15,0.5))
```

```
plot(rw4.l$x, type = "l")
abline(h=c(0.15,0.5))
```



Det må sies å være til forveksling likt i konvergens i tillegg konvergerer $\sigma = 0.05$ noe fortere.

Problem 2.3


```
set.seed(123)

#initialize constants and parameters
N <- 5000 #length of chain
burn <- 1000 #burn-in length
X <- matrix(0, N, 2) #the chain, a bivariate sample
rho <- -.9 #correlation
mu_x <- 0
mu_y <- 0
sigma_x <- 1
sigma_y <- 1
s1 <- sqrt(1-rho^2)*sigma_x
s2 <- sqrt(1-rho^2)*sigma_y

##### generate the chain #####
X[1, ] <- c(mu_x, mu_y) #initialize
for (i in 2:N) {
  x2 <- X[i-1, 2]
  m_x <- mu_x + rho * (y - mu_y) * sigma_x/sigma_y
  X[i, 1] <- rnorm(1, m_x, s1)
  x_x <- X[i, 1]
  m_y <- mu_y + rho * (x_x - mu_x) * sigma_y/sigma_x
  X[i, 2] <- rnorm(1, m_y, s2)
}
b <- burn + 1
x <- X[b:N, ]

hist(sampled_ray_dist, breaks="scott",
     main="Gibbs Sampler for N(0,1)", xlab="x", freq=FALSE)
```

Gibbs Sampler for $N(0,1)$



Task 3

problem 3.1


```
p1 <- 0.5
n = 10
data <- p1 * rnorm(n, 1,1) + (1-p1) * rnorm(n, 3,1)
data1 <- p1 * rgamma(n, rate = 1, shape = 1) + (1-p1) * rgamma(n, rate = 3, shape = 1)
par(mfrow = c(1, 1))
print(density(data))
```

```
##
## Call:
## density.default(x = data)
##
## Data: data (10 obs.); Bandwidth 'bw' = 0.4101
##
##      x              y
## Min.   :-0.01707   Min.   :0.001771
## 1st Qu.: 1.05666   1st Qu.:0.053027
## Median : 2.13038   Median :0.283447
## Mean   : 2.13038   Mean    :0.232491
## 3rd Qu.: 3.20411   3rd Qu.:0.374453
## Max.   : 4.27783   Max.    :0.455501
```


```
df <- approxfun(density(data)) ##extract the density function##
xnew <- seq(0,10,0.2)
n <- length(data)
h1 <- 0.9 * min(c(IQR(data)/1.34, sd(data))) * n^(-1/5) #density seems not unimodal#
h2 <- 1.06 * sd(data) * n^(-1/5)
```

lets first look at the extream example (high, low):


```
par(mfrow = c(2, 2))
plot(density(data,bw=.10) , main="extream narrow bw - gaus distr") # "Day of the tentacle"
plot(density(data,bw=1000), main="extream wide bw - gaus distr") #Glææt - "too much
plot(density(data1,bw=.10), main="extream narrow bw - gamma distr") # "Day of the tentacle"
plot(density(data1,bw=1000), main="extream narrow bw - gamma distr") #Glææt - "too much
```

extream narrow bw – gaus distr


N = 10 Bandwidth = 0.1

extream wide bw – gaus distr

N = 10 Bandwidth = 1000

extream narrow bw – gamma distr

N = 10 Bandwidth = 0.1


extream narrow bw – gamma distr

N = 10 Bandwidth = 1000

From this we can conclude that choise of bandwith is essential. To narrow you get spikes, to wide you obtain a very flat, smooth curve converged towards the normal distribution. As we also see the result is the same if we change the kernel distribution. So lets try to fit a better curve:

```
#par(mfrow=c(2,2))
plot(density(data)) #The default method applies the Gaussian kernel#3
points(xnew,df(xnew),col=2) #bulky
```



density.default(x = data)



N = 10 Bandwidth = 0.4101

```
plot(density(data,bw=h1)) #Better, this include the duality of the two distr.
```


density.default(x = data, bw = h1)



N = 10 Bandwidth = 0.4101

```
plot(density(data,bw=h2)) #oversmooth...
```

density.default(x = data, bw = h2)



N = 10 Bandwidth = 0.483

problem 3.2

we have:

$$\begin{aligned}\mathbf{x} &= \{x_1, \dots, x_n\} \sim \text{gamma}(\theta) \\ \theta &= (r, \lambda)\end{aligned}$$

we first need to solve for the MLE functions for the gamma distributed RVs:

$$\begin{aligned}f(\mathbf{x}|\theta) &= \frac{\lambda^r}{\Gamma(r)} \cdot \mathbf{x}^{r-1} \cdot \exp(-\lambda \mathbf{x}) \\ L(\theta|x_1, \dots, x_n) &= \prod_{i=1}^n f(\mathbf{x}|\theta) \\ &= \frac{\lambda^{n \cdot r}}{\Gamma(r)^n} \cdot \sum_i^n x_i^{r-1} \cdot \exp(-\lambda \sum_i^n x_i)\end{aligned}$$

we obtain the log likelihood by applying magic

$$l(\theta|x_1, \dots, x_n) = nr \cdot \log(\lambda) - n \cdot \log(\Gamma(r)) + (r-1) \cdot \sum_i^n \log(x_i) - \lambda \sum_i^n x_i$$

we then take the derivative to find the maximum:

$$\begin{aligned}\frac{\partial l}{\partial \lambda} &= \frac{nr}{\lambda} - \sum_{i=1}^n x_i = 0 \\ \frac{\partial l}{\partial r} &= n \cdot \log(\lambda) - \frac{n\Gamma'(r)}{\Gamma(r)} + \sum_{i=1}^n \log(x_i) = 0\end{aligned}$$

to find $\hat{\lambda}$ is fairly straight forward:

$$\begin{aligned}\frac{nr}{\lambda} - \sum_{i=1}^n x_i &= 0 \\ \frac{nr}{\lambda} &= \sum_{i=1}^n x_i \\ \hat{\lambda} &= \frac{nr}{\sum_{i=1}^n x_i} = \frac{r}{n^{-1} \sum_{i=1}^n x_i} = \frac{\hat{r}}{\bar{x}_n}\end{aligned}$$

to find \hat{r} is more of a pain since we have the $\frac{\Gamma'(r)}{\Gamma(r)}$ term also known as digamma:

$$\psi(r) = \frac{\Gamma'(r)}{\Gamma(r)} = \frac{\partial}{\partial r} \log(\Gamma(r))$$

to simplify we substitute $\hat{\lambda}$ into the equation:

$$n \cdot \log\left(\frac{\hat{r}}{\bar{x}_n}\right) - \frac{n\Gamma'(\hat{r})}{\Gamma(\hat{r})} + \sum_{i=1}^n \log(x_i) = 0 \quad (1)$$

$$\log\left(\frac{\hat{r}}{\bar{x}_n}\right) - \frac{1}{n} \sum_{i=1}^n \log(x_i) = \frac{\Gamma'(\hat{r})}{\Gamma(\hat{r})} = \psi(\lambda \bar{x}_n) \quad (2)$$

for $\bar{x}_n = \frac{r}{\lambda}$

Due to the digamma function the MLE needs to be solved numerically by root solver (in R `uniroot`), which is presumably a bisection method, since derivative is not needed. Bisection only needs a continuous function to converge to one of the existing roots. Below the same simulation as from “Statistical computing with R” by Maria L. Rizzo” page 340 is conducted (with comments on the process)

```
m <- 20000
est <- matrix(0, m, 2)
n <- 200
r <- 5
lambda <- 2

# equation (1) above, of which we want to find the root
obj <- function(lambda, xbar, logx.bar) {
  digamma(lambda * xbar) - logx.bar - log(lambda)
}

# Do m simulations of a set of  $x_1, \dots, x_n$ 
for (i in 1:m) {
  x <- rgamma(n, shape=r, rate=lambda)
  xbar <- mean(x)

  # using uniroot to solve numerically
  u <- uniroot(obj, lower = .001, upper = 10e5,
    xbar = xbar, logx.bar = mean(log(x)))

  # columns vector of MLE simulations
  lambda.hat <- u$root
  r.hat <- xbar * lambda.hat
  est[i, ] <- c(r.hat, lambda.hat)
}

# mean of the MLE simulations
ML <- colMeans(est)

ML

## [1] 5.070865 2.030602
```